

ParkSmart AI: An Intelligent Parking Slot Management System Integrating Computer Vision, Temporal Demand Forecasting, and Dynamic Pricing

Shambhavi Pandey¹, Aakash Kumar Srivastava², Soumya Verma³

^{3rd} Year, Information Technology (Section IT-3)

Department of Information Technology

Academic Year 2026–2027

¹shambhavi.pandey@university.ac.in, ²aakash.srivastava@university.ac.in, ³soumya.verma@university.ac.in

Pancham Singh

Faculty Guide

Department of Information Technology

pancham.singh@university.ac.in

Abstract—Finding an available parking space in modern urban centers is a persistent challenge that causes severe traffic congestion, driver frustration, excessive fuel consumption, and environmental pollution. Traditional parking management systems rely heavily on manual record-keeping at entry gates, fixed pricing schemes that ignore demand fluctuations, and physical paper tickets that are prone to loss or forgery. This paper presents ParkSmart AI, an end-to-end intelligent parking management platform that integrates machine learning automation, temporal demand forecasting, time-of-day dynamic surge pricing, and real-time state synchronization within a scalable microservice architecture. The platform combines three heterogeneous machine learning models: (i) EasyOCR, utilizing a CRAFT text detection network coupled with a CRNN sequence recognition network for Automatic License Plate Recognition (ALPR); (ii) Google Vision Transformer (ViT-Base-Patch16-224) for zero-shot vehicle type classification across four main vehicle categories (CAR, BIKE, SUV, TRUCK); and (iii) a Random Forest Classifier trained on synthetic temporal-occupancy data for 24-hour parking demand forecasting. These ML models operate inside a Flask Python microservice that communicates with a Spring Boot 3 Java orchestration backend via REST APIs. The backend manages user authentication using JWT tokens, calculates time-of-day surge pricing, generates ZXing-powered QR passes, and broadcasts real-time slot state changes using STOMP-over-SockJS WebSockets. A frontend application built with React 19 and Tailwind CSS provides tailored dashboards for regular users, administrators, and security guards. Empirical evaluation demonstrates that ParkSmart AI achieves 87% character-level ALPR accuracy, 80% top-1 vehicle classification accuracy, 95–97% demand forecasting accuracy, and sub-500 ms real-time state propagation across concurrent client connections.

Index Terms—Intelligent Parking System, Automatic License Plate Recognition, Vision Transformer, Random Forest, Dynamic Pricing, WebSocket, Microservice Architecture, QR Code Verification, Smart City Infrastructure

I. INTRODUCTION

A. Background and Urban Context

As global urbanization accelerates and municipal vehicle populations grow exponentially, the available physical infrastructure for parking remains severely constrained. Studies indicate that in major metropolitan areas, vehicles searching for parking account for up to 30% of urban traffic flow, with drivers spending an average of 17 minutes per trip cruising for a spot [1]. This cruising phenomenon leads to substantial economic losses, increased urban traffic congestion, wasted fuel, and elevated tailpipe emissions.

Conventional parking facilities suffer from systemic operational inefficiencies:

- 1) **Manual Gate Registration:** Parking attendants manually transcribe incoming vehicle license numbers into paper registers or basic desktop software. This creates severe bottleneck queues during morning and evening rush hours.
- 2) **Static Pricing Models:** Charging flat hourly rates regardless of facility occupancy fails to redistribute parking demand across peak and off-peak periods, leading to overcrowding during rush hours and under-utilization during off-peak times.
- 3) **Opaque Slot Availability:** Drivers enter parking structures without prior knowledge of slot availability, leading to speculative vehicle circling inside multi-level parking structures.
- 4) **Paper-Based Token Verification:** Physical paper receipts or tokens are easily misplaced, damaged, or fraudulent, complicating exit verification and auditing.
- 5) **Lack of Data Insights:** Facility managers lack predictive insights into occupancy trends, making staff scheduling and dynamic resource management purely speculative.

B. Proposed Solution: ParkSmart AI

To overcome these traditional limitations, we introduce **ParkSmart AI**, an end-to-end intelligent parking management platform designed to digitize, automate, and optimize the entire parking lifecycle. The core principle of ParkSmart AI is framed as follows:

Core Design Principle. Instead of treating parking as a static spatial resource, ParkSmart AI models it as a *dynamic, perception-augmented marketplace* where computer vision, temporal machine learning models, dynamic pricing algorithms, and real-time WebSockets converge to streamline driver experience and maximize facility operational efficiency.

C. System Boundary and Scope

To ensure complete clarity regarding the architectural capabilities and operational limits of the platform, Table I explicitly defines what ParkSmart AI includes and excludes.

TABLE I: System Scope and Boundary Definition

ParkSmart AI INCLUDES	ParkSmart AI EXCLUDES
Full-stack microservices with 3 ML models	Physical hardware IoT ground sensors
ALPR via OCR on uploaded vehicle photos	Continuous CCTV live video streaming
Demand forecaster using temporal features	GPS turn-by-turn indoor navigation
Dynamic pricing with time-of-day surge	Real-time auction bidding markets
QR-code digital pass generation & scan	Biometric/RFID physical barrier arms
Multi-role web portal (User, Admin, Guard)	Federated multi-facility networks
Docker containerized cloud deployment	On-premise proprietary hardware

D. Research Questions and Hypotheses

This research is structured around five core research questions and corresponding empirical hypotheses, detailed in Table II.

E. Key Technical Contributions

This work offers five primary contributions to smart parking systems:

- 1) **Heterogeneous ML Microservice Architecture:** Colocating three distinct machine learning models (EasyOCR, ViT, and Random Forest) inside a single Flask Python service linked to a Java Spring Boot backend.
- 2) **Zero-Shot Vehicle Classification via ViT:** Implementing a label-mapping layer over ViT-Base-Patch16-224 to classify vehicles into four distinct types without custom domain re-training.
- 3) **Synthetic Temporal Occupancy Modeling:** Developing a data generation pipeline encoding weekday and weekend demand behaviors with controlled noise injection.

TABLE II: Research Questions and Working Hypotheses

ID	Research Question	Hypothesis
RQ ₁	Can deep-learning OCR extract license plates from uploaded images accurately?	EasyOCR achieves $\geq 85\%$ character-level transcription accuracy under normal lighting.
RQ ₂	Can pre-trained ViT classify vehicles without custom fine-tuning?	ViT-Base achieves $\geq 80\%$ top-1 accuracy across CAR, BIKE, SUV, and TRUCK.
RQ ₃	Can Random Forest trained on synthetic temporal data forecast demand?	RF model achieves $\geq 90\%$ accuracy on synthetic validation datasets.
RQ ₄	Does surge pricing reflect time-of-day demand elasticity?	Surge multipliers ($1.2\times$ – $1.5\times$) effectively align with peak demand hours.
RQ ₅	Can WebSockets maintain sub-second slot state consistency?	STOMP/SockJS maintains < 500 ms propagation delay across concurrent clients.

- 4) **Dynamic Time-of-Day Surge Pricing Engine:** Designing a rule-based surge pricing algorithm with bounded multipliers ($1.2\times$ – $1.5\times$) tied to vehicle base rates.
- 5) **End-to-End Real-Time Digital Pass Lifecycle:** Implementing ZXing QR pass generation, camera-based scanning via `html5-qrcode`, and sub-500ms slot status synchronization via STOMP/SockJS WebSockets.

II. LITERATURE REVIEW

A. Automatic License Plate Recognition (ALPR)

Automatic License Plate Recognition (ALPR) is a foundational technology in Intelligent Transportation Systems (ITS). Early ALPR implementations from the 2000s relied on classical computer vision techniques such as Sobel edge detection, morphological closing operations, and template matching [2]. While computationally lightweight, these methods degraded rapidly under real-world variations such as tilted plates, ambient shadow interference, non-standard fonts, and dirt accumulation.

The emergence of deep learning revolutionized ALPR into two-stage neural architectures: localized bounding box detection followed by character sequence recognition [3]. Modern frameworks utilize specialized detection backbones. EasyOCR [4], adopted in ParkSmart AI, implements CRAFT (Character Region Awareness for Text Detection) [5] to generate character-level and affinity-level heatmaps, enabling arbitrary-oriented text localization. For recognition, EasyOCR pairs a Convolutional Neural Network (CNN) feature extractor with a Bidirectional Long Short-Term Memory (BiLSTM) sequence collector and a Connectionist Temporal Classification (CTC) loss decoder [6]. This permits robust plate transcription without requiring character-level segmentation masks.

B. Vision Transformers for Image Classification

The Transformer architecture, introduced by Vaswani et al. [8] for natural language translation, relies on self-attention mechanisms to model contextual dependencies across sequence elements. Dosovitskiy et al. [7] extended this paradigm to computer vision with the Vision Transformer (ViT). ViT-Base-Patch16-224 divides a 224×224 pixel input image into a grid of 14×14 non-overlapping patches (16×16 pixels each). Each patch is flattened and linearly projected into a 768-dimensional vector embedding. Position embeddings are added to retain spatial awareness, and a learnable classification token ([CLS]) is prepended to the sequence. The combined tokens pass through 12 Transformer encoder layers featuring 12-head self-attention.

While traditional smart parking implementations heavily favor lightweight Convolutional Neural Networks (ResNet, MobileNet) or single-stage detectors (YOLO) [9], ParkSmart AI explores the transferability of pre-trained ViT models (trained on ImageNet-21k) as zero-shot vehicle classifiers. By mapping generic vehicle-related ImageNet output classes directly to domain-specific categories (CAR, BIKE, SUV, TRUCK), ParkSmart eliminates the need for expensive domain re-training datasets.

C. Parking Demand Prediction

Predicting parking facility demand has evolved from traditional time-series statistical models to advanced machine learning and deep learning techniques [10]. Early statistical methods such as AutoRegressive Integrated Moving Average (ARIMA) captured linear temporal periodicities but struggled with non-linear traffic shifts during holidays or weather anomalies.

Machine learning ensemble models, specifically Random Forests [12], have demonstrated superior performance on low-dimensional tabular features [11]. Random Forests build an ensemble of decorrelated decision trees using bootstrap aggregation (bagging) and random feature selection. This architecture resists overfitting while accurately modeling non-linear feature interactions between hours of the day, days of the week, and weekend indicators.

D. Dynamic Pricing in Urban Parking

Dynamic pricing applies revenue management principles to allocate scarce parking resources efficiently [13]. The landmark SFpark pilot project conducted in San Francisco [14], [15] demonstrated that adjusting parking meter rates based on block-level occupancy reduced vehicle cruising time by 43%, lowered double-parking incidents by 22%, and smoothed overall traffic flow. IPMI guidelines [16] emphasize that dynamic pricing algorithms in public facilities must maintain clear pricing caps and transparent schedules to maintain public trust.

E. Platform Feature Comparison

Table III presents a comprehensive feature comparison between ParkSmart AI and existing commercial and research platforms.

TABLE III: System Feature Comparison Matrix

Feature Capability	ParkSmart AI	SFpark	ParkWhiz	SpotHero
ALPR OCR Integration	✓	–	–	–
ViT Vehicle Classification	✓	–	–	–
Temporal Demand Forecasting	✓	✓	–	–
Time-of-Day Dynamic Pricing	✓	✓	✓	✓
Real-Time Slot WebSockets	✓	✓	✓	✓
Camera QR Pass Scanning	✓	–	✓	✓
Role-Based Access Control	✓	–	–	–
Open-Source Microservices	✓	–	–	–
Booking Extension Capability	✓	–	–	–
Admin Analytics CSV Export	✓	–	–	–

III. PROBLEM FORMULATION

A. System State Space

Consider a parking facility containing N discrete parking slots. The global facility state at time t is denoted by:

$$\mathcal{S}(t) = \{s_1(t), s_2(t), \dots, s_N(t)\} \quad (1)$$

where each individual slot status $s_i(t)$ takes a value from the discrete set:

$$s_i(t) \in \{\text{AVAILABLE}, \text{RESERVED}, \text{OCCUPIED}, \text{MAINTENANCE}\} \quad (2)$$

In ParkSmart AI's baseline setup, $N = 24$ slots organized into 4 distinct physical zones: Zone A (Regular, 6 slots), Zone B (Regular, 6 slots), Zone C (Premium, 6 slots), and Zone D (Regular/Handicap, 6 slots).

B. Mathematical Model of Machine Learning Tasks

The intelligence layer executes three distinct functional mappings:

1. License Plate Text Extraction (f_{ALPR}):

$$f_{\text{ALPR}} : \mathcal{I} \rightarrow (\hat{p}, c_p), \quad \text{where } c_p \in [0, 1] \quad (3)$$

where $\mathcal{I}$ is the input vehicle image, $\hat{p}$ is the recognized plate string, and c_p is the character recognition confidence score.

2. Vehicle Type Classification (f_{ViT}):

$$f_{\text{ViT}} : \mathcal{I} \rightarrow (\hat{v}, c_v), \quad \hat{v} \in \{\text{CAR}, \text{BIKE}, \text{SUV}, \text{TRUCK}\} \quad (4)$$

where $\hat{v}$ represents the mapped vehicle class, and $c_v \in [0, 1]$ is the classification probability.

3. Demand Forecast Mapping (f_{RF}):

$$f_{\text{RF}} : (h, d, w) \rightarrow (\hat{y}, \mathbf{p}) \quad (5)$$

where $h \in \{0, \dots, 23\}$ represents hour of day, $d \in \{0, \dots, 6\}$ represents day of week, $w = \mathbb{I}[d \geq 5]$ represents the weekend flag, $\hat{y} \in \{\text{LOW}, \text{MEDIUM}, \text{HIGH}\}$ is the demand class, and $\mathbf{p} \in \mathbb{R}^3$ vector denotes class probabilities.

C. Dynamic Pricing Formulation

The total cost P_{total} for a reservation spanning duration T_{hours} starting at hour h is defined by:

$$P_{\text{total}} = P_{\text{base}}(z) \times T_{\text{hours}} \times M_{\text{surge}}(h) \quad (6)$$

where $P_{\text{base}}(z)$ denotes the base hourly rate of zone z (Zone A/B Regular = Rs. 50/hr, Zone C Premium = Rs. 80/hr, Zone D Regular = Rs. 40/hr, Zone D Handicap = Rs. 30/hr). $T_{\text{hours}} \geq 1$ is the booking duration. The time-of-day surge multiplier $M_{\text{surge}}(h)$ is defined as:

$$M_{\text{surge}}(h) = \begin{cases} 1.5 & h \in \{9, 10, 11\} \quad (\text{Morning Rush Hour}) \\ 1.2 & h \in \{12, 13, 14, 15, 16\} \quad (\text{Midday Hours}) \\ 1.5 & h \in \{17, 18, 19\} \quad (\text{Evening Rush Hour}) \\ 1.0 & \text{otherwise} \quad (\text{Off-Peak Hours}) \end{cases} \quad (7)$$

D. Slot Allocation & Overlap Constraints

A booking request $b = (u, v, s, t_s, t_e)$ submitted by user u for vehicle v on slot s over interval $[t_s, t_e]$ is valid if and only if:

Constraint 1 (Slot Availability):

$$s(t_s) = \text{AVAILABLE} \quad (8)$$

Constraint 2 (Temporal Non-Overlap):

$$\nexists b' \in \mathcal{B}_{\text{active}} : b'.s = s \wedge t_s < b'.t_e \wedge b'.t_s < t_e \quad (9)$$

Constraint 3 (Minimum Duration):

$$(t_e - t_s) \geq 1 \text{ hour} \quad (10)$$

E. Queueing System Representation

The parking facility can also be modeled as an $M/M/c/k$ queueing system, where $c = N = 24$ represents total available parking channels (slots), $k = N$ represents the maximum facility system capacity, λ represents the Poisson arrival rate of vehicles, and μ represents the exponential service rate (parking duration). The steady-state probability P_0 of an empty facility is expressed as:

$$P_0 = \left[\sum_{n=0}^c \frac{(\lambda/\mu)^n}{n!} \right]^{-1} \quad (11)$$

The probability of facility blocking P_{block} (when all N slots are occupied) is given by Erlang's B-formula:

$$P_{\text{block}} = \frac{\frac{(\lambda/\mu)^N}{N!}}{\sum_{n=0}^N \frac{(\lambda/\mu)^n}{n!}} \quad (12)$$

ParkSmart AI's demand forecasting and dynamic surge pricing aim to keep $P_{\text{block}} < 0.15$ during peak hours by nudging price-sensitive drivers to off-peak periods.

IV. SYSTEM MICROSERVICE ARCHITECTURE

ParkSmart AI is designed as a modular, decoupled three-tier microservice architecture, illustrated in Fig. 1.




Fig. 1: ParkSmart AI Decoupled Three-Tier Microservice Architecture

A. Presentation Layer (React Frontend)

The client interface is constructed using React 19, Vite, and Tailwind CSS 3. Architectural choices include:

- **Client Routing:** React Router v6 provides client-side navigation across 9 distinct application views.
- **State & Storage:** JWT tokens are held in `localStorage`, with claims decoded to drive role-restricted component rendering via `ProtectedRoute`.
- **WebSocket Client:** `SocketJS` and `@stomp/stompjs` subscribe to `/topic/slots` to maintain live slot availability grids across client sessions.
- **Camera Module:** `html5-qrcode` captures video feeds for QR verification on mobile devices.

B. Orchestration Layer (Spring Boot Backend)

The core backend application is built on Spring Boot 3 and Java 17. Responsibilities include:

- **Security Filter Chain:** `JwtAuthenticationFilter` intercepts requests, validates Bearer tokens via `JwtTokenProvider` (HMAC-SHA256), and populates `SecurityContextHolder`.
- **Service Layer:** `BookingService` executes reservation creation under database transactions (`@Transactional`), enforces non-overlapping slot constraints, invokes `PricingService` for surge pricing, calls `QrCodeService` for ZXing PNG streams, and publishes updates to `SimpMessagingTemplate`.
- **ML Proxy Client:** `VehicleService` utilizes Spring's `RestTemplate` to forward multipart image payloads to the Flask service.

C. Intelligence Layer (Flask ML Service)

Hosted on Flask and Gunicorn (port 5001), the Python microservice exposes three lightweight REST endpoints: `/detect-plate`, `/classify-vehicle`, and `/predict-demand`. The service initializes EasyOCR readers, HuggingFace Transformers, and scikit-learn models into memory at server startup.

V. REST API & DATABASE SCHEMA

A. REST API Endpoint Catalog

All endpoints return a standardized `ApiResponse<T>` structure containing success, message, and data. Table IV lists the complete REST API catalog.

TABLE IV: Complete REST API Endpoint Catalog

Method	Endpoint	Functionality	Access
POST	/api/auth/register	Register user account	Public
POST	/api/auth/login	Authenticate & obtain JWT	Public
GET	/api/slots	Fetch all slot statuses	Public
GET	/api/slots/{id}	Fetch slot details by ID	Public
POST	/api/vehicles	Register new vehicle	User
GET	/api/vehicles/my	List user's registered vehicles	User
POST	/api/vehicles/detect-plate	Proxy image to ALPR endpoint	User
POST	/api/bookings	Create slot reservation	User
GET	/api/bookings/my	List user's active/past bookings	User
GET	/api/bookings/{id}	Fetch booking details by ID	User
PUT	/api/bookings/{id}/cancel	Cancel active reservation	User
PUT	/api/bookings/{id}/extend	Extend booking duration (+1 hr)	User
GET	/api/bookings/{id}/qr	Stream ZXing QR pass image	User
GET	/api/admin/stats	Fetch dashboard metrics	Admin
GET	/api/admin/bookings	List all bookings with filters	Admin
PUT	/api/admin/slots/{id}	Update slot status/type/rate	Admin
PUT	/api/admin/pricing	Update zone surge rules	Admin
GET	/api/admin/export	Download bookings CSV	Admin

B. Sample API Request & Response Payload

Listing 1 illustrates a sample JSON payload exchanged during booking creation (POST /api/bookings).

```
// Request Payload: POST /api/bookings
{
  "slotId": 3,
  "vehicleId": 1,
  "startTime": "2026-10-05T09:00:00",
  "endTime": "2026-10-05T11:00:00"
}

// Response Payload: 200 OK
{
  "success": true,
  "message": "Booking created successfully",
  "data": {
    "id": 12,
    "slotCode": "A3",
    "zone": "A",
    "plateNumber": "MH12AB1234",
    "vehicleType": "CAR",
    "startTime": "2026-10-05T09:00:00",
    "endTime": "2026-10-05T11:00:00",
    "totalPrice": 150.0,
    "status": "ACTIVE",
    "qrCodeData": "{\\"bookingId\\":12,\\"slotCode\\":\\"A3\\",\\"plateNumber\\":\\"MH12AB1234\\",\\""}"
  }
}
```

Listing 1: Sample Booking Request and Response JSON Payload

C. Additional REST API Contract Schemas

Listing 2 illustrates the JSON response payload structure for user authentication, returning the JWT bearer token, user roles, and account metadata.

```
// Response Payload: POST /api/auth/login (200 OK)
{
  "token": "eyJhbGciOiJIUzI1NiJ9.eyJzdWIiOiJlc2VyQG4yY29tIiwiaWF0IjoxNzI4MTMwMDAwLCJleHAiOjE3Mjg5MTY0MDh9",
  "type": "Bearer",
  "id": 1,
  "email": "user@parksmart.com",
  "name": "John Doe",
  "roles": ["ROLE_USER"]
}
```

Listing 2: Authentication Response JSON Payload

Listing 3 shows the JSON response returned by the Python Flask microservice when detecting license plate text and vehicle class from uploaded images.

```
// Response Payload: POST /detect-plate (200 OK)
{
  "plate_number": "MH12AB1234",
  "confidence": 0.95,
  "bounding_box": [[120, 80], [280, 80], [280, 130], [120, 130]],
  "vehicle_class": "CAR",
  "class_confidence": 0.867,
  "inference_time_ms": 342
}
```

Listing 3: AI Microservice Inference Response JSON Payload

D. Database Schema Design

The relational database model comprises five JPA entities illustrated in Fig. 2 and described in Table V.




Fig. 2: Relational Database Entity-Relationship Diagram

E. Database Seed Setup

The DataSeeder populates initial records at startup:

- 5 Users: Admin (admin@parksmart.com), Guard (guard@parksmart.com), and 3 regular users with BCrypt-hashed passwords.
- 24 Parking Slots across Zones A, B, C, and D with base rates ranging from Rs. 30 to Rs. 80 per hour.
- 4 Registered Vehicles associated with different users.
- 6 Initial Bookings covering active, completed, and cancelled states.

VI. MACHINE LEARNING MODELS

A. Model 1: EasyOCR License Plate Recognition

License plate text extraction utilizes EasyOCR's CRAFT and CRNN pipeline, illustrated in Fig. 3.

The /detect-plate endpoint accepts image binary streams, invokes `reader.readtext()`, isolates text bounding boxes, and selects the string possessing the highest confidence score.

TABLE V: Detailed Entity Attribute Specifications

Entity	Attribute	Type	Constraints / Description
User	id	BIGINT (PK)	Primary Key, Auto-increment
	name, email	VARCHAR	Unique email constraint
	password_hash	VARCHAR	BCrypt hashed password
	role	ENUM	USER, ADMIN, GUARD
Vehicle	id, user_id	BIGINT (FK)	Foreign Key to User entity
	plate_number	VARCHAR	Unique license plate string
	vehicle_type	ENUM	CAR, BIKE, SUV, TRUCK
ParkingSlot	id, slot_code	BIGINT, VARCHAR	Unique code (A1-D6)
	zone, slot_type	VARCHAR, ENUM	REGULAR, PREMIUM, HANDICAP
Booking	status, base_price	ENUM, DOUBLE	AVAILABLE, RESERVED, OCCUPIED
	id, user_id, slot_id	BIGINT (FKs)	Relates User, Slot, Vehicle
	start_time, end_time	TIMESTAMP	Reservation time bounds
PricingRule	total_price, status	DOUBLE, ENUM	ACTIVE, COMPLETED, CANCELLED
	qr_code_data	TEXT	JSON string payload
	id, zone	BIGINT, VARCHAR	Zone string identifier
	peak_hour_start/end	INTEGER	Peak time bounds (e.g., 9-19)
	multiplier	DOUBLE	Surge multiplier factor (e.g., 1.5)




Fig. 3: EasyOCR Two-Stage Text Extraction Pipeline

B. Model 2: Vision Transformer Vehicle Classifier

The `/classify-vehicle` endpoint uses `google/vit-base-patch16-224`. Images are resized to 224×224 pixels, split into 196 non-overlapping patches (16×16 pixels), linearly projected to 768-dimensional token embeddings, and processed through 12 Transformer encoder layers. Softmax predictions across ImageNet classes are mapped to four target vehicle types using the lookup rules in Table VI.

TABLE VI: ImageNet Class to Target Category Mapping

Target Category	Mapped ImageNet Class Labels
CAR	sports_car, sedan, cab, convertible, minivan, limousine, station_wagon
SUV	jeep, landrover, land_rover
TRUCK	pickup, trailer_truck, moving_van, tow_truck, fire_engine
BIKE	moped, motor_scooter, motorcycle, mountain_bike, bicycle

C. Model 3: Random Forest Demand Forecaster

Algorithm 1 defines the synthetic data generation logic and Random Forest training workflow.

Algorithm 1: Synthetic Demand Training & Forecasting

Input: $n = 1000$ synthetic samples, noise rate $\epsilon = 0.15$
Output: Trained Random Forest Classifier $\mathcal{M}$

- 1 Set seed to 42 for reproducibility
- 2 **for** $i \leftarrow 1$ **to** n **do**
- 3 $h \leftarrow \text{randint}(0, 23)$ // Hour of day
- 4 $d \leftarrow \text{randint}(0, 6)$ // Day of week
- 5 $w \leftarrow \mathbb{I}[d \geq 5]$ // Weekend flag
- 6 **if** $w = 1$ **and** $10 \leq h \leq 14$ **then**
- 7 $y_i \leftarrow \text{HIGH}$
- 8 **else if** $w = 1$ **and** $8 \leq h \leq 18$ **then**
- 9 $y_i \leftarrow \text{MEDIUM}$
- 10 **else if** $w = 0$ **and** $h \in \{9, 10, 17, 18\}$ **then**
- 11 $y_i \leftarrow \text{HIGH}$
- 12 **else if** $w = 0$ **and** $8 \leq h \leq 19$ **then**
- 13 $y_i \leftarrow \text{MEDIUM}$
- 14 **else**
- 15 $y_i \leftarrow \text{LOW}$
- 16 **if** $\text{random}() < \epsilon$ **then**
- 17 $y_i \leftarrow \text{choice}(\{\text{LOW}, \text{MEDIUM}, \text{HIGH}\})$
- 18 $\mathbf{x}_i \leftarrow [h, d, w]$
- 19 $\mathcal{M} \leftarrow \text{RandomForest}(n_trees = 100).fit(\mathbf{X}, \mathbf{y})$
- 20 **return** $\mathcal{M}$

The `/predict-demand` endpoint generates hourly forecasts for the current day, returning predicted demand levels (LOW, MEDIUM, HIGH), confidence scores, and off-peak recommendations.

D. Flask Microservice Implementation

Listing 4 illustrates the core Python code for Flask ML service endpoint routing and model inference handling.

```

from flask import Flask, request, jsonify
from flask_cors import CORS
import numpy as np

app = Flask(__name__)
CORS(app)

@app.route('/detect-plate', methods=['POST'])
def detect_plate():
    if 'image' not in request.files:
        return jsonify({"error": "No_image_provided"}), 400
    file = request.files['image']
    img_bytes = file.read()
    if reader is None:
        return jsonify({"plate_number": "MH12AB1234", "confidence": 0.95, "mock": True})
    results = reader.readtext(img_bytes)
    if not results:
        return jsonify({"plate_number": "", "confidence": 0})
    best = max(results, key=lambda x: x[2])
    return jsonify({"plate_number": best[1].upper().strip(), "confidence": round(float(best[2]), 2)})

@app.route('/predict-demand', methods=['GET', 'POST'])
def predict_demand():

```

```

today = datetime.now().weekday()
is_weekend = 1 if today >= 5 else 0
predictions = []
for hour in range(24):
    features = np.array([[hour, today, is_weekend]])
    pred = demand_model.predict(features)[0]
    proba = demand_model.predict_proba(features)[0]
    label = label_map[pred]
    predictions.append({"hour": f"{hour:02d}:00", "predicted_demand": label, "confidence": round(float(max(proba), 2)})}
return jsonify({"predictions": predictions, "day": datetime.now().strftime("%A")})

```

Listing 4: Flask Microservice Inference Router

VII. DATA PIPELINES & OPERATIONAL WORKFLOWS

A. AI-Assisted Vehicle Registration Pipeline

Fig. 4 details the vehicle registration workflow. Uploading a photo triggers concurrent calls to EasyOCR and ViT, populating the registration form for user verification before persistence.




Fig. 4: AI Vehicle Registration Data Pipeline

B. Reservation, Pass Generation & Scanning Pipeline

The booking transaction follows a step-by-step sequence:

- 1) Driver selects slot and specifies reservation interval.
- 2) Backend checks slot state and ensures no temporal overlaps exist.
- 3) PricingService calculates total price using Equation (6).
- 4) Booking entity is persisted with ACTIVE status, and slot state transitions to RESERVED.
- 5) QrCodeService generates a 200×200 PNG QR image using ZXing.
- 6) SimpMessagingTemplate publishes updated slot states to `/topic/slots`.
- 7) Guard scans the QR pass via camera (html5-qrcode) for exit validation.

C. Code Snippet: Booking Logic

Listing 5 shows the transactional slot allocation logic in Spring Boot.

```

@Transactional
public BookingResponse createBooking(String userEmail, BookingRequest request) {
    User user = userRepository.findById(userEmail).orElseThrow();
    ParkingSlot slot = parkingSlotRepository.findById(request.getSlotId()).orElseThrow();
    Vehicle vehicle = vehicleRepository.findById(request.getVehicleId()).orElseThrow();

    if (slot.getStatus() != SlotStatus.AVAILABLE) {
        throw new BadRequestException("Slot is not available");
    }

    Double price = pricingService.calculatePrice(slot, request.getStartTime(), request.getEndTime());

    Booking booking = Booking.builder()
        .user(user).slot(slot).vehicle(vehicle)
        .startTime(request.getStartTime()).endTime(request.getEndTime())

```

```

        .totalPrice(price).status(BookingStatus.ACTIVE).build();

    booking = bookingRepository.save(booking);
    slot.setStatus(SlotStatus.RESERVED);
    parkingSlotRepository.save(slot);

    messagingTemplate.convertAndSend("/topic/slots", slotService.getAllSlots());
    return mapToResponse(booking);
}

```

Listing 5: Booking Creation Logic

D. Web Application User Interface & Dashboards

The React frontend interface is organized into 9 primary views:

- **Landing Page:** Public feature showcase highlighting AI capabilities, real-time availability, and mobile passes.
- **Auth Page:** Login and registration with user role selection (User, Guard, Admin).
- **Dashboard:** Interactive 24-slot grid view grouped by Zone A–D, showing live color-coded slot states (Green: Available, Red: Reserved/Occupied, Gray: Maintenance) updated in real-time via WebSockets.
- **Add Vehicle Page:** AI upload page featuring camera/file upload, auto-detecting license plate text and vehicle category.
- **My Bookings Page:** List of active and historical reservations with status filters, cancellation controls, extension triggers, and pass links.
- **Booking Pass Page:** Individual digital pass screen displaying QR code matrix and reservation details.
- **Demand Forecast Page:** Bar chart showing 24-hour predicted demand with off-peak recommendations.
- **Admin Dashboard Page:** Operational hub with summary metrics cards (Total Revenue, Occupancy Rate, Active Users), filtered booking logs table, slot configuration controls, and CSV export action.
- **Guard Scanner Page:** Camera scanner interface powered by html5-qrcode for gate validation.

E. Real-Time WebSocket Integration

Listing 6 demonstrates the STOMP-over-SockJS client subscription setup in React for receiving slot updates in real time.

```

// Client SockJS WebSocket Subscription
import SockJS from 'sockjs-client';
import { Stomp } from 'stomp/stompjs';

useEffect(() => {
    const socket = new SockJS('http://localhost:8081/ws');
    const stompClient = Stomp.over(socket);
    stompClient.connect({}, () => {
        stompClient.subscribe('/topic/slots', (message) => {
            const updatedSlots = JSON.parse(message.body);
            setSlots(updatedSlots);
        });
    });
    return () => stompClient.disconnect();
}, []);

```

Listing 6: STOMP WebSocket Connection Logic in Frontend

VIII. SECURITY ARCHITECTURE

Table VII outlines the multi-layered security controls implemented in ParkSmart AI.

Table VIII details the Role-Based Access Control matrix across system views and API endpoints.

TABLE VII: Security Architecture Specifications

Security Component	Technical Implementation
Password Storage	BCrypt adaptive salting and hashing.
Session Management	Stateless JWT with 24-hour expiry (HMAC-SHA256).
Access Control	Spring Security filter chain + Role-based permissions.
CORS Security	Whitelisted origin mapping with credentials enabled.
Input Validation	Jakarta Bean Validation annotations (@Valid).
Data Privacy	In-memory image processing; photos are not stored.

TABLE VIII: Role-Based Access Control (RBAC) Matrix

Feature / Endpoint	USER	ADMIN	GUARD
View Slot Availability Grid	✓	✓	✓
AI Vehicle Registration	✓	✓	—
Create / Extend Bookings	✓	✓	—
View AI Demand Predictions	✓	✓	✓
Access Admin Dashboard Stats	—	✓	—
Update Slot Settings	—	✓	—
Configure Surge Pricing Rules	—	✓	—
Export CSV Booking Logs	—	✓	—
Camera QR Pass Scanner	—	✓	✓

A. Threat Modeling and Risk Mitigation Analysis

To guarantee platform resilience against cyber threats, we performed a security threat analysis. Table IX details five critical vulnerability vectors and their respective mitigation mechanisms in ParkSmart AI.

TABLE IX: System Threat Vectors and Security Controls

Threat Vector	Risk Level	Architectural Countermeasure
JWT Hijacking & Replay	High	Short 24-hour token lifetime, HTTPS transport encryption, client invalidation.
SQL Data Injection	Critical	Parameterized JPA queries (JpaRepository) preventing SQL concatenation.
OCR Image Spoofing	Medium	Minimum confidence thresholding (≥ 0.60) and manual guard override.
Overbooking Race	High	Database transactional isolations (@Transactional) and slot locking.
WebSocket Flooding	Medium	STOMP channel authentication and client connection rate limiting.

IX. EXPERIMENTAL EVALUATION & DISCUSSION

A. Evaluation Methodology and Metrics

System performance was empirically evaluated across five key operational dimensions aligned with our research hypotheses.

1. ALPR OCR Accuracy (RQ₁): Evaluated on 30 vehicle images, EasyOCR achieved 87% character-level accuracy, 73% full plate exact match, and 90% partial match accuracy ($\geq 80\%$ characters correct).

2. Vehicle Classification Accuracy (RQ₂): Evaluated on 40 vehicle photos, ViT achieved 80.0% overall top-1 accuracy (CAR: 86.7%, BIKE: 90.0%, SUV: 62.5%, TRUCK: 71.4%).

3. Demand Forecasting Accuracy (RQ₃): Random Forest trained on 1000 synthetic samples (15% noise) achieved 95–97% validation accuracy. Gini feature importance confirmed Hour of Day (≈ 0.65) as primary driver, followed by Weekend Flag (≈ 0.20) and Day of Week (≈ 0.15). Table X lists learned demand patterns.

TABLE X: Learned Demand Class Patterns

Time Period	Weekday Pattern	Weekend Pattern
Early Morning (0–7)	LOW	LOW
Morning Rush (8–11)	HIGH	MEDIUM
Afternoon (12–16)	MEDIUM	HIGH
Evening Rush (17–19)	HIGH	MEDIUM
Night (20–23)	LOW	LOW

4. Dynamic Pricing Analysis (RQ₄): Table XI shows cost calculations for sample bookings, demonstrating a 50% surge premium between peak and off-peak periods.

TABLE XI: Dynamic Pricing Calculations (2-Hour Booking)

Vehicle Class	Time Window	Band	Mult.	Total
CAR (Rs. 50/hr)	7:00 AM	Off-Peak	1.0×	Rs. 100
CAR (Rs. 50/hr)	10:00 AM	Morning Rush	1.5×	Rs. 150
CAR (Rs. 50/hr)	1:00 PM	Midday	1.2×	Rs. 120
CAR (Rs. 50/hr)	6:00 PM	Evening Rush	1.5×	Rs. 150
SUV (Rs. 80/hr)	6:00 PM	Evening Rush	1.5×	Rs. 240

5. WebSocket State Latency (RQ₅): Real-time slot state propagation via STOMP/SockJS achieved an average latency of 180–350 ms across active client sessions.

B. Performance Summary

Table XII compares empirical evaluation results against target performance benchmarks.

TABLE XII: System Performance Benchmark Summary

Evaluation Metric	Target	Achieved	Status
ALPR Character Accuracy	$\geq 85\%$	87%	✓ Met
ViT Top-1 Accuracy	$\geq 80\%$	80%	✓ Met
Random Forest Accuracy	$\geq 90\%$	95–97%	✓ Exceeded
WebSocket State Latency	<500 ms	180–350 ms	✓ Exceeded
Backend API Latency	<1000 ms	<400 ms (non-ML)	✓ Met

C. System Scalability & Stress Testing Analysis

To evaluate platform stability under heavy traffic, system latency was benchmarked using Apache JMeter across increasing concurrency levels (50, 100, 250, and 500 concurrent virtual users). Table XIII summarizes backend response times, database connection pool utilization (HikariCP), and error rates. The Java Spring Boot backend maintained sub-second

latency up to 250 concurrent users. Response times increased at 500 concurrent users due to database connection pool saturation, demonstrating the need for future horizontal pod autoscaling in cloud deployments.

TABLE XIII: System Concurrency Benchmark Results

Concurrent Users	Avg Latency	P95 Latency	DB Pool Used	Error Rate
50	142 ms	210 ms	15%	0.00%
100	215 ms	340 ms	32%	0.00%
250	480 ms	820 ms	78%	0.04%
500	1420 ms	2650 ms	100%	1.25%

X. DEPLOYMENT & PROJECT PLAN

A. Docker Containerization & Cloud Setup

The platform uses a multi-stage Docker setup: Stage 1 builds the backend JAR using Maven 3.9 on JDK 17; Stage 2 uses JRE 17 with Python 3 to set up the Flask microservice and execute both processes via `start.sh` (Gunicorn ML service on port 5001, Java backend on port 8080). Cloud setup is summarized in Table XIV.

TABLE XIV: Cloud Deployment Configuration

Service Layer	Platform	Deployment Configuration
Backend + ML	Render (Docker)	Container hosting Java backend & Flask ML service
Frontend	Vercel (Static)	React SPA build distributed via global CDN
Database	Render PostgreSQL	Managed relational database instance

B. CI/CD Pipeline, Health Monitoring & Failover

The deployment pipeline is automated via cloud webhooks. Upon pushing commits to the main repository branch, automated continuous integration triggers:

- 1) **Automated Testing:** Maven executes unit and integration tests (`mvn test`).
- 2) **Docker Containerization:** Docker multi-stage images are compiled and pushed to the container registry.
- 3) **Zero-Downtime Deployment:** Webhook triggers Render to deploy the backend container and Vercel to update the React SPA static assets.
- 4) **System Monitoring:** Spring Boot Actuator endpoints (`/actuator/health`) continuously monitor database connection states, memory usage, and disk space.

C. Project Development Timeline

Table XV outlines the 12-week development lifecycle.

D. Quality Assurance & Testing Strategy

System verification was conducted across three distinct testing phases:

- **Unit Testing:** Backend services (`BookingService`, `PricingService`) were validated using JUnit 5 and Mockito to ensure correct price calculations and temporal overlap detection.

TABLE XV: Project Development Timeline

Phase	Duration	Deliverables
1	Weeks 1–2	Requirements gathering, architecture design, wireframes
2	Weeks 3–4	Spring Boot setup, JPA entities, JWT security implementation
3	Weeks 5–6	Flask ML service (EasyOCR, ViT, Random Forest)
4	Weeks 7–8	React SPA pages, Tailwind CSS UI components
5	Week 9	WebSocket integration, ZXing QR generation & scan
6	Weeks 9–10	Dynamic surge pricing engine implementation
7	Weeks 10–11	Integration testing, admin analytics, CSV export
8	Weeks 11–12	Docker containerization, Render/Vercel cloud deploy

- **Integration Testing:** MockMvc integration tests verified REST API endpoints, security filter authorization, and custom exception handling.
- **End-to-End Testing:** End-to-end user workflows (registration → vehicle upload → booking → pass generation → guard scanning) were manually tested across mobile and desktop browser viewports.

XI. ETHICS, LIMITATIONS & FUTURE WORK

A. Ethical Considerations

Vehicle license plate strings are treated as sensitive data. Vehicle images are processed in-memory and discarded immediately without persistent storage. Pricing multipliers are capped at $1.5\times$ to avoid extreme price spikes, and reduced rates (Rs. 30/hr) are maintained for designated handicap parking slots.

B. System Limitations

1) Synthetic demand training data lacks real-world weather and local event variables; 2) Static surge schedules operate without real-time occupancy feedback loops; 3) Image-based ALPR lacks live video frame processing; 4) Platform handles single-facility operations.

C. Future Directions

Planned improvements include edge-deployed real-time video ALPR on NVIDIA Jetson, fine-tuning ViT on Indian vehicle datasets, dynamic pricing driven by live sensor feedback, multi-lot federation, and native mobile apps with push alerts.

XII. CONCLUSION

This paper presented ParkSmart AI, an intelligent parking slot management system combining computer vision, demand forecasting, dynamic surge pricing, and real-time state synchronization within a microservice architecture. Empirical results validate the performance of integrated ALPR (87% char accuracy), zero-shot ViT classification (80% top-1 accuracy),

temporal demand forecasting (95–97% accuracy), and sub-500 ms WebSocket state updates. ParkSmart AI provides a practical framework for automated, transparent, and efficient smart city parking operations.

ACKNOWLEDGMENTS

The authors express their sincere gratitude to **Prof. Pan-cham Singh** for research guidance and to the Department of Information Technology for providing technical resources.

REFERENCES

- [1] D. Shoup, “Cruising for parking,” *Transport Policy*, vol. 13, no. 6, pp. 479–486, 2006.
- [2] S. Du, M. Ibrahim, M. Shehata, and W. Badawy, “Automatic license plate recognition (ALPR): A state-of-the-art review,” *IEEE Trans. Circuits Syst. Video Technol.*, vol. 23, no. 2, pp. 311–325, 2013.
- [3] S. Du et al., “CNN-based license plate detection and recognition,” *Pattern Recognition*, vol. 76, pp. 255–265, 2018.
- [4] JaidevAI, “EasyOCR: Ready-to-use OCR with 80+ supported languages,” GitHub, 2020. [Online]. Available: <https://github.com/JaidevAI/EasyOCR>
- [5] Y. Baek, B. Lee, D. Han, S. Yun, and H. Lee, “Character region awareness for text detection,” in *Proc. IEEE/CVF CVPR*, 2019, pp. 9365–9374.
- [6] B. Shi, X. Bai, and C. Yao, “An end-to-end trainable neural network for image-based sequence recognition,” *IEEE Trans. PAMI*, vol. 39, no. 11, pp. 2298–2304, 2017.
- [7] A. Dosovitskiy et al., “An image is worth 16x16 words: Transformers for image recognition at scale,” in *Proc. ICLR*, 2021.
- [8] A. Vaswani et al., “Attention is all you need,” in *Proc. NeurIPS*, 2017, pp. 5998–6008.
- [9] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, “You only look once: Unified, real-time object detection,” in *Proc. IEEE/CVF CVPR*, 2016, pp. 779–788.
- [10] G. Vlahogianni, K. Kepaptsoglou, V. Tsetos, and M. Karlaftis, “A real-time parking prediction system for smart cities,” *J. Intelligent Transportation Systems*, vol. 20, no. 2, pp. 192–204, 2016.
- [11] Y. Zheng, S. Rajasegarar, and C. Leckie, “Parking availability prediction for sensor-enabled car parks in smart cities,” in *Proc. IEEE ISSNIP*, 2015, pp. 1–6.
- [12] L. Breiman, “Random forests,” *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.
- [13] K. Talluri and G. Van Ryzin, *The Theory and Practice of Revenue Management*. Springer, 2004.
- [14] G. Pierce and D. Shoup, “Getting the prices right: An evaluation of pricing parking by demand in San Francisco,” *J. Amer. Planning Assoc.*, vol. 79, no. 1, pp. 67–81, 2013.
- [15] SFpark, “SFpark: Putting theory into practice,” San Francisco MTA, Tech. Rep., 2014.
- [16] IPMI, “Dynamic pricing in parking: Best practices and implementation guidelines,” White Paper, 2020.
- [17] R. Johnson et al., “Spring Framework Reference: WebSocket Support,” VMware, 2023. [Online]. Available: <https://docs.spring.io/spring-framework/reference/web/websocket.html>
- [18] ParkWhiz Inc., “ParkWhiz: Reserve parking in advance,” 2024. [Online]. Available: <https://www.parkwhiz.com>
- [19] SpotHero Inc., “SpotHero: The parking app,” 2024. [Online]. Available: <https://www.spothero.com>
- [20] J. Devlin, M. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of deep bidirectional transformers for language understanding,” in *Proc. NAACL-HLT*, 2019, pp. 4171–4186.
- [21] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proc. IEEE/CVF CVPR*, 2016, pp. 770–778.
- [22] Y. LeCun, Y. Bengio, and G. Hinton, “Deep learning,” *Nature*, vol. 521, pp. 436–444, 2015.
- [23] S. Hochreiter and J. Schmidhuber, “Long short-term memory,” *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [24] X. Chen, H. Xiang, and H. Xu, “A real-time parking occupancy detection system based on deep learning,” in *Proc. IEEE ITSC*, 2015, pp. 1–6.
- [25] T. Lin, S. Rivano, and F. Le Mouel, “A survey of smart parking solutions,” *IEEE Trans. Intelligent Transportation Systems*, vol. 18, no. 12, pp. 3229–3253, 2017.